

TOPIC 6/60

What is the Attention Mechanism?

The secret engine inside the Transformer.

How AI learned to focus on what matters.

Ashish Jangra

[linkedin.com/in/ashish-jangra](https://www.linkedin.com/in/ashish-jangra)



THE BOTTLENECK

Before Attention, AI had a massive "information bottleneck" problem when translating or reading text.

The Old Way (Seq2Seq):

It tried to compress an entire 50-word sentence into a single, fixed-size mathematical vector before translating it.


> Result: It forgot the beginning of the sentence by the time it reached the end.

Ashish Jangra

[linkedin.com/in/ashish-jangra](https://www.linkedin.com/in/ashish-jangra)



THE HUMAN WAY



How do we read?

When translating a sentence, humans don't memorize the whole thing as a single block. We **pay attention** to specific words.

"The boy grabbed his **umbrella** because it was **raining**."

To understand "umbrella", we naturally focus on the word "raining", ignoring words like "the" or "his".

Ashish Jangra

[linkedin.com/in/ashish-jangra](https://www.linkedin.com/in/ashish-jangra)



THE CORE CONCEPT

Self-Attention

Letting words talk

Self-Attention gives every word in a sentence the ability to look at **every other word**, and calculate how strongly they are connected.

It asks:

"I am the word 'Apple'. Out of all the other words in this sentence, which ones should I care about to understand my own meaning?"

Ashish Jangra

[linkedin.com/in/ashish-jangra](https://www.linkedin.com/in/ashish-jangra)



THE ARCHITECTURE

Query, Key, Value

The Filing Cabinet Analogy

To calculate attention, the AI assigns three roles (vectors) to every single word.

Query (Q): What I am looking for (The search bar).

Key (K): What I contain (The video title on YouTube).

Value (V): My actual meaning (The video itself).

Ashish Jangra

[linkedin.com/in/ashish-jangra](https://www.linkedin.com/in/ashish-jangra)



HOW IT CALCULATES

The Score

Matrix Multiplication

The AI takes a word's **Query**, and multiplies it against every other word's **Key**.

The Output:

It produces a percentage score (from 0.0 to 1.0). High score = High relevance.

> It then multiplies this score by the "Value" to update the word's final context.

Ashish Jangra

[linkedin.com/in/ashish-jangra](https://www.linkedin.com/in/ashish-jangra)



DYNAMIC MEANING

SENTENCE A

"I deposited money in the bank."

"Bank" pays high attention to "money" and "deposited".

SENTENCE B

"I sat on the bank of the river."

"Bank" pays high attention to "river" and "sat". Its mathematical vector completely changes.

Ashish Jangra

[linkedin.com/in/ashish-jangra](https://www.linkedin.com/in/ashish-jangra)



TAKING IT FURTHER

Multi-Head

Why just look at the sentence one way? Transformers use "Multi-Head Attention" to run this process multiple times simultaneously.

Head 1 might pay attention to **grammar** (who is the subject?).

Head 2 might pay attention to **emotion** (is this positive or negative?).

Ashish Jangra

[linkedin.com/in/ashish-jangra](https://www.linkedin.com/in/ashish-jangra)



WHY IT'S REVOLUTIONARY

No Bottlenecks

It doesn't compress text. Every word has direct access to every other word, instantly.

Long-Range Context

A word on page 10 can perfectly "attend" to a word on page 1, solving the short-term memory problem.

Highly Parallel

Because all matrices are calculated at once, it utilizes GPU power with extreme efficiency.

Ashish Jangra

[linkedin.com/in/ashish-jangra](https://www.linkedin.com/in/ashish-jangra)



Attention Mechanism

CONCEPT UNLOCKED



Query: What am I looking for?



Key: What do I contain?



Value: The actual meaning



Connects all words dynamically

CONNECT WITH ME




LinkedIn



GitHub



Kaggle



Instagram

Ashish Jangra

[linkedin.com/in/ashish-jangra](https://www.linkedin.com/in/ashish-jangra)

