

TOPIC 18/60

What is Top-P in LLMs?

Also known as Nucleus Sampling.

The intelligent way to filter AI choices.

Ashish Jangra

linkedin.com/in/ashish-jangra



THE RIGIDITY OF TOP-K

Top-K always picks from a fixed number of words. But sometimes the model's confidence varies wildly.

Context: "The cat is..."

Small (90%), cute (8%), sleeping (1%)

Here, K=50 would include 47 terrible words that don't make sense.

Ashish Jangra

[linkedin.com/in/ashish-jangra](https://www.linkedin.com/in/ashish-jangra)



TOP-P DEFINED



Nucleus Sampling

Instead of a fixed *number* of words (K), Top-P uses a **probability threshold** (P).

If $P=0.9$: The AI considers only the smallest set of words whose probabilities **add up to 90%**. It ignores the long tail of low-probability noise.

Ashish Jangra

[linkedin.com/in/ashish-jangra](https://www.linkedin.com/in/ashish-jangra)



THE ADVANTAGE

Adaptability

Top-P is dynamic. It adjusts the size of the vocabulary pool based on the AI's confidence.

High Confidence?

Top-P might only consider the top 2 words (P=90%).

Low Confidence?

Top-P might automatically expand to consider 200 words (P=90%).

Ashish Jangra

[linkedin.com/in/ashish-jangra](https://www.linkedin.com/in/ashish-jangra)



THE "NUCLEUS"

The Top Tail

Filtering the Junk

Think of the probability distribution as an atom. The "Nucleus" (the top probability mass) contains the high-quality choices.

Top-P cuts off the "electrons" (low-probability tail) where nonsense lives, keeping only the meaningful core.

Ashish Jangra

[linkedin.com/in/ashish-jangra](https://www.linkedin.com/in/ashish-jangra)



THE SETTINGS

Top-P vs. Top-K

Most modern LLM interfaces let you toggle these, but they serve the same goal: avoiding gibberish.

Summary:

K = Fixed amount (e.g., top 50 choices).

> P = Fixed quality (e.g., top 90% of probability mass).

Ashish Jangra

[linkedin.com/in/ashish-jangra](https://www.linkedin.com/in/ashish-jangra)



WHY USE TOP-P?

WITHOUT TOP-P

Model picks from a rigid bucket.

Sometimes the bucket is too big (too much noise),
sometimes it's too small (not enough variety).

WITH TOP-P

Model breathes with the context.

It maintains coherence across varying confidence levels,
leading to more natural output.

Ashish Jangra

[linkedin.com/in/ashish-jangra](https://www.linkedin.com/in/ashish-jangra)



THE DOWNSIDE

Over-Smoothing

If you set P too low (e.g., 0.1), the model becomes boring and repetitive, as it only picks the most certain words.

Most balanced setting is $P=0.9$. It keeps the model creative while pruning the catastrophic nonsense of the long tail.

Ashish Jangra

linkedin.com/in/ashish-jangra



WHY IT MATTERS

Natural Flow

Top-P produces text that feels "human," balancing logic with creative variance.

Reliability

It significantly reduces the risk of "runaway" text generation where the AI starts babbling.

Standard Practice

Almost all state-of-the-art LLMs (GPT-4, etc.) use a form of Top-P (Nucleus) sampling by default.

Ashish Jangra

[linkedin.com/in/ashish-jangra](https://www.linkedin.com/in/ashish-jangra)



