

TOPIC 27/60

What is LoRA?

Low-Rank Adaptation:
The fine-tuning revolution.

Ashish Jangra

[linkedin.com/in/ashish-jangra](https://www.linkedin.com/in/ashish-jangra)



THE FINE-TUNING BOTTLENECK

Full fine-tuning requires updating every single weight in a 70B+ model. That requires massive GPU clusters and millions of dollars.

The LoRA Insight:

Weights in a massive model are "over-parameterized." Most of them don't need to change for a specific task. We only need to adjust a tiny fraction.

Ashish Jangra

[linkedin.com/in/ashish-jangra](https://www.linkedin.com/in/ashish-jangra)



THE LORA METHOD



Freezing the Core

Instead of training the giant weight matrix, we freeze it completely. We then inject tiny "adapter" layers alongside it.

During training, only those tiny adapter layers are updated. The main model never changes.

Ashish Jangra

[linkedin.com/in/ashish-jangra](https://www.linkedin.com/in/ashish-jangra)



THE MATH

Low-Rank Matrices

Instead of training a giant matrix (e.g., 4096×4096), we decompose it into two tiny matrices (e.g., 4096×8 and 8×4096).

The Impact:

This reduces the number of trainable parameters by **over 99.9%**. You can fine-tune a massive LLM with just a fraction of the VRAM.

Ashish Jangra

[linkedin.com/in/ashish-jangra](https://www.linkedin.com/in/ashish-jangra)



FLEXIBILITY

Modular AI

Because LoRA adapters are just small files (often a few MBs), you can "hot-swap" your model's personality.

Use the same base LLaMA model, but swap LoRA adapters to turn it into a **Legal Expert**, a **Medical Advisor**, or a **Poetry Writer** in seconds.

Ashish Jangra

[linkedin.com/in/ashish-jangra](https://www.linkedin.com/in/ashish-jangra)



WHY LORA WINS

Efficiency

- **Hardware:** Run fine-tuning on consumer-grade GPUs (e.g., RTX 3090/4090).
- **Portability:** Adapters are tiny; share them on HuggingFace easily.
- **Performance:** LoRA achieves nearly identical results to full fine-tuning.

Ashish Jangra

[linkedin.com/in/ashish-jangra](https://www.linkedin.com/in/ashish-jangra)



THE ECOSYSTEM

QLORA

Quantized LoRA

Combines Quantization (4-bit) with LoRA. Allows fine-tuning of 70B+ models on extremely limited home hardware.

MERGING

Zero-Latency

You can merge LoRA weights back into the base model, so there's zero performance penalty during inference.

Ashish Jangra

[linkedin.com/in/ashish-jangra](https://www.linkedin.com/in/ashish-jangra)



THE TRADE-OFF

Rank Selection

The "Rank" (R) is the hyperparameter you choose. A higher rank captures more complex behavior but requires more VRAM.

Most people use Rank 8 or 16. Going too high defeats the purpose of LoRA; going too low may not capture the nuances of the task.

Ashish Jangra

[linkedin.com/in/ashish-jangra](https://www.linkedin.com/in/ashish-jangra)



SUMMARY OF IMPACT

Full Fine-Tuning

Expensive, huge hardware requirement, slow.

LoRA Fine-Tuning

Cheap, accessible, fast, modular, identical performance.

Ashish Jangra

[linkedin.com/in/ashish-jangra](https://www.linkedin.com/in/ashish-jangra)



LoRA

CONCEPT UNLOCKED



Freezes base weights



Trains only tiny adapters



Saves VRAM & time



Modular and swappable

CONNECT WITH ME



LinkedIn



GitHub



Kaggle



Instagram

Ashish Jangra

[linkedin.com/in/ashish-jangra](https://www.linkedin.com/in/ashish-jangra)

