

TOPIC 29/60

What is LoRA?

Low-Rank Adaptation.

The mathematical hack that democratized open-source AI.

Ashish Jangra

linkedin.com/in/ashish-jangra



THE 70 BILLION PROBLEM

Fine-tuning a massive model like Llama 3 (70B parameters) traditionally requires updating all 70 billion numbers.

The Reality:

This demands terabytes of VRAM and clusters of \$30,000 GPUs. It made fine-tuning impossible for 99% of developers and researchers.

Ashish Jangra

[linkedin.com/in/ashish-jangra](https://www.linkedin.com/in/ashish-jangra)



THE INSIGHT



Intrinsic Dimension

In 2021, Microsoft researchers realized something profound: When you fine-tune an AI for a specific task, you don't actually need to change all its knowledge.

The "changes" needed are actually very simple. They exist in a **low-rank subspace**. You only need to update a tiny fraction of the neural network's logic.

Ashish Jangra

[linkedin.com/in/ashish-jangra](https://www.linkedin.com/in/ashish-jangra)



THE ANALOGY

Textbooks & Notes

Think of a base LLM as a massive, heavy, 10,000-page encyclopedia.

Full Fine-Tuning:

Rewriting the entire encyclopedia from scratch just to update a few facts.

LoRA:

Sticking a small **Post-it note** on a page with your updates. You leave the heavy book exactly as it is.

Ashish Jangra

[linkedin.com/in/ashish-jangra](https://www.linkedin.com/in/ashish-jangra)



UNDER THE HOOD

The Matrix Hack

Neural networks are just massive grids of numbers (Matrices). Instead of updating a giant 10,000 x 10,000 matrix...

$W + (A \times B)$

W: The original, massive frozen matrix.

A & B: Two tiny new matrices (e.g., 10,000 x 8 and 8 x 10,000).

Ashish Jangra

[linkedin.com/in/ashish-jangra](https://www.linkedin.com/in/ashish-jangra)



THE IMPACT

Storage Miracle

Because we are only training the tiny "A" and "B" matrices, the output of our training isn't a massive 140GB model file.

140 GB → 50 MB

The resulting LoRA "adapter" file is usually under 100 Megabytes. It's so small you can email it or attach it to a Discord message.

Ashish Jangra

[linkedin.com/in/ashish-jangra](https://www.linkedin.com/in/ashish-jangra)



HARDWARE IMPACT

BEFORE LORA

Enterprise Clusters Only

Required 8x A100 GPUs (\$120,000+) to fine-tune a decent sized open-source model.

AFTER LORA

Consumer Hardware

You can now fine-tune a 7B or 8B model on a single consumer RTX 3090 or 4090 gaming GPU in a few hours.

Ashish Jangra

[linkedin.com/in/ashish-jangra](https://www.linkedin.com/in/ashish-jangra)



SUPERPOWER

Hot Swapping

Because the adapters are just tiny files applied on top of the frozen base model, you can instantly swap an AI's personality.

Load the Base Llama Model.

- + Load `Coding_LoRA.safetensors` → (AI writes Python)
- + Swap to `Medical_LoRA.safetensors` → (AI reads charts)

Ashish Jangra

[linkedin.com/in/ashish-jangra](https://www.linkedin.com/in/ashish-jangra)



THE TRADE-OFFS

Slight Performance Drop

A LoRA is generally 95% to 98% as good as a full fine-tune. For most, the massive cost savings are worth the tiny accuracy hit.

New Knowledge is Hard

LoRA is amazing for teaching an AI a new *style* or *format*. It struggles to inject massive amounts of brand new factual knowledge.

Ashish Jangra

[linkedin.com/in/ashish-jangra](https://www.linkedin.com/in/ashish-jangra)



LoRA

CONCEPT UNLOCKED

-  Freezes base model weights
-  Trains tiny A/B matrices
-  Democratized AI training
-  Enables instant hot-swapping

CONNECT WITH ME



LinkedIn



GitHub



Kaggle



Instagram

Ashish Jangra

linkedin.com/in/ashish-jangra

