

TOPIC 36/60

LLMs vs SLMs

The battle of Parametric Scale.

Do you need a trillion-parameter giant running in a massive server farm, or a highly-optimized expert running directly on your laptop? Let's explore the divide.

Ashish Jangra

linkedin.com/in/ashish-jangra



THE TAXONOMY

Defining the Scale

The defining metric between a "Large" and "Small" language model is its parameter count—the number of neural connections that store its knowledge and logic.



LLM (Large Language Model)

Generally 70 Billion to 1 Trillion+ Parameters

Requires massive cloud GPU clusters to run. Exhibits high-level reasoning and broad world knowledge. (e.g., GPT-4, Claude 3.5 Opus).



SLM (Small Language Model)

Generally Under 10 Billion Parameters

Can run locally on edge devices (MacBooks, phones). Highly optimized for specific tasks like summarization or code generation. (e.g., Phi-3, Llama 3 8B).

Ashish Jangra

[linkedin.com/in/ashish-jangra](https://www.linkedin.com/in/ashish-jangra)



THE GENERALISTS


The Case for LLMs

Frontier models are built to be universal reasoning engines. They contain vast encyclopedic knowledge and can adapt to nearly any task zero-shot.



Emergent Reasoning

At massive scale, models develop advanced logical deduction, coding abilities, and math skills without explicit fine-tuning.



World Knowledge

They have memorized vast amounts of the internet, acting as intelligent knowledge bases across a multitude of domains.



The Downside

Massive latency (slower to generate words), incredibly high API costs, and data must leave your network to reach a cloud provider.

Ashish Jangra

[linkedin.com/in/ashish-jangra](https://www.linkedin.com/in/ashish-jangra)



THE AGILE EXPERTS

The Case for SLMs

If you only need an AI to extract dates from receipts, or summarize emails, using a 1-Trillion parameter model is like driving an 18-wheeler to pick up groceries.



Blazing Fast Latency

Fewer parameters mean less math per token. SLMs can generate text almost instantly, perfect for real-time applications.



Air-Gapped Privacy

Because they fit in normal RAM, SLMs can run entirely offline on local enterprise servers. Your data never leaves your building.



The Trade-off

SLMs have limited world knowledge. If you ask a 3B parameter model for a niche historical fact, it is much more likely to hallucinate.

Ashish Jangra

[linkedin.com/in/ashish-jangra](https://www.linkedin.com/in/ashish-jangra)



THE RESEARCH BREAKTHROUGH

Textbooks Are All You Need

For years, AI labs believed the only way to make models smarter was to make them *bigger*. Microsoft challenged this with the **Phi** models, proving that Data Quality > Data Quantity.

The Synthetic Data Revolution

Instead of training a small model on the noisy, messy, unfiltered public internet (Reddit, broken code, spam), researchers used GPT-4 to generate millions of "textbook quality" examples of logic, math, and coding.

The Result:

A tiny 3-Billion parameter model (Phi-2) trained on highly curated textbook data began outperforming 70-Billion parameter models trained on internet garbage.

Ashish Jangra

[linkedin.com/in/ashish-jangra](https://www.linkedin.com/in/ashish-jangra)




DEPLOYMENT REALITIES

Hardware Economics

The choice between an LLM and an SLM often comes down to IT infrastructure and operational budgets rather than pure intelligence.


70B+ Parameter LLMs

Llama 3 70B, Qwen 72B



Massive VRAM Required

Requires ~40GB to 80GB+ of VRAM just to load into memory.




High Compute Cost

Requires expensive clusters of Nvidia A100 or H100 GPUs (Thousands of dollars/month).

Edge Native


< 8B Parameter SLMs

Llama 3 8B, Phi-3 Mini



Low Memory Footprint

Quantized to 4-bit, they use only 4GB to 6GB of standard RAM.



Consumer Hardware

Runs blazingly fast on standard MacBooks, iPhones, or cheap CPU servers.

Ashish Jangra

[linkedin.com/in/ashish-jangra](https://www.linkedin.com/in/ashish-jangra)



THE ENTERPRISE PLAY

Fine-Tuning Specialists

While an SLM lacks general world knowledge out-of-the-box, it is incredibly easy and cheap to **fine-tune** it into a master of one specific domain.

The "Intern to Specialist" Pipeline



Base SLM (8B)

Okay at everything,
master of nothing.

FINE-TUNE
(\$50
COMPUTE)



Medical SLM

Trained on 10,000 hospital discharge notes. Now extracts patient symptoms with 99% accuracy, beating zero-shot GPT-4.

Ashish Jangra


[linkedin.com/in/ashish-jangra](https://www.linkedin.com/in/ashish-jangra)



SYSTEM ARCHITECTURE

LLM + SLM Routing

Modern AI applications don't choose just one model. They use a **Router** to direct traffic based on the complexity of the user's prompt, saving massive amounts of money and time.



SUMMARY GUIDE

When to use which?

REQUIREMENT	GIANT LLMS (70B+)	EDGE SLMS (<8B)
Latency / Speed	Slower (Cloud round-trip)	Near-Instant
Operating Cost	High (API pay-per-token)	Free / Electricity only
Data Privacy	Data sent to OpenAI/Anthropic	100% Air-gapped / Local
Complex Reasoning	Excellent (System 2 Thinking)	Struggles with deep logic

Ashish Jangra

[linkedin.com/in/ashish-jangra](https://www.linkedin.com/in/ashish-jangra)



The Future is Hybrid.

The era of passing every tiny query to a trillion-parameter cloud model is ending.

Tomorrow's architecture relies on local SLMs acting as fast, private agents, escalating to cloud LLMs only when deep intellect is required.

CONNECT WITH ME



LinkedIn



GitHub



Kaggle



Instagram

Ashish Jangra

[linkedin.com/in/ashish-jangra](https://www.linkedin.com/in/ashish-jangra)

