

TOPIC 47/60

Metadata Filtering in RAG

Precision Beyond Semantics.

How structured attributes keep vector similarity searches highly relevant, fully compliant, and completely contextual.

Ashish Jangra

linkedin.com/in/ashish-jangra



THE PROBLEM

The Blindness of Space

Vector databases are amazingly good at finding conceptual matches. However, they are mathematically blind to structured constraints.

The Semantic Leak

If you ask: *"How do I configure API keys in v2.0 of our product?"*

A pure vector search might retrieve the highest similarity chunk from the v1.0 documentation because the prose is linguistically identical. It lacks the logic to know 1.0 != 2.0.

Ashish Jangra

[linkedin.com/in/ashish-jangra](https://www.linkedin.com/in/ashish-jangra)



THE SOLUTION

Structured Meta-Tags



What is Metadata Filtering?

It is the process of appending structured key-value attributes directly to your text chunks inside the vector database.

Dynamic Payload Attributes:

- Date: 2026-07-16
- Region: EMEA
- Permission: Level_3_Admin
- Version: v2.0

Ashish Jangra

[linkedin.com/in/ashish-jangra](https://www.linkedin.com/in/ashish-jangra)



THE PIPELINE CHALLENGE

Naive Filtering Methods

Historically, developers combined metadata and vectors using two disjoint steps:

1. Post-Filtering (Similarity First)

Query Top-100 vectors, then throw out everything that doesn't match the metadata filter.

⚠️ **Issue:** If all Top-100 are from v1.0, you end up with exactly 0 results.

2. Pre-Filtering (Metadata First)

Filter metadata using SQL, then perform vector search on the remaining subset.

⚠️ **Issue:** Performs poorly on indexes because it forces linear scanning over millions of points.

Ashish Jangra

[linkedin.com/in/ashish-jangra](https://www.linkedin.com/in/ashish-jangra)



MODERN ENGINE ARCHITECTURE

Single-Stage Filtering

Modern vector databases (Qdrant, Pinecone, Milvus) perform vector traversal and boolean filter resolution ****simultaneously****.

```
// Simultaneous Graph Traversal & Pruning
```

```
Step 1: Landing on Index Graph Node (Coordinate Similarity: 0.89)
```

```
Step 2: Asserting Metadata Boolean filter: { "version": "v2.0" }
```

```
✅ MATCH: Retain node in the traversal queue.
```

```
Step 3: Traversing edge to Node 2 (Coordinate Similarity: 0.94)
```

```
Step 4: Asserting Metadata Boolean: { "version": "v1.0" }
```

```
❌ NO-MATCH: Prune graph branch instantly. Do not evaluate descendants.
```

```
* Yields logarithmic  $O(\log N)$  lookup speeds without sacrificing semantic recall.
```

Ashish Jangra

[linkedin.com/in/ashish-jangra](https://www.linkedin.com/in/ashish-jangra)



SECURITY & COMPLIANCE

Hard Multi-Tenancy

In enterprise RAG, preventing data leaks across different corporate clients is the absolute highest priority.

The Security Threat

If client documents are stored together in a single vector index, a conceptual query from Tenant A might bleed and retrieve Tenant B's sensitive financial data.

The Metadata Isolation Shield

By prepending `tenant_id: tenant_902` to every database query, the database filters the index dynamically, guaranteeing absolute data isolation with zero risk of cross-tenant exposure.

Ashish Jangra

[linkedin.com/in/ashish-jangra](https://www.linkedin.com/in/ashish-jangra)



QUERY ANATOMY

Payload Syntax

What a multi-dimensional query payload looks like in a production vector database:

```
{
  "vector": [0.124, -0.982, 0.455, ...],
  "top_k": 5,
  "filter": {
    "must": [
      { "key": "department", "match": "Finance" },
      { "key": "year", "gt": 2024 }
    ],
    "must_not": [
      { "key": "classification", "match": "Internal_Draft" }
    ]
  }
}
```

Ashish Jangra

[linkedin.com/in/ashish-jangra](https://www.linkedin.com/in/ashish-jangra)



RELATIONAL HYBRIDIZATION

SQL & Spatial Logic

Nested Logical Operations

Modern search allows boolean operators (AND, OR, NOT, IN) directly nested inside spatial vector ranges, matching complex business rules perfectly.

Distance Threshold Filters

Only return documents that pass structured filters *AND* exceed a cosine similarity threshold of 0.82, preventing the system from outputting low-confidence responses.

Ashish Jangra

[linkedin.com/in/ashish-jangra](https://www.linkedin.com/in/ashish-jangra)



METADATA TAXONOMY

Structural vs. Semantic

How we split and populate metadata payloads inside our document processing pipeline:

Structural Meta

Determined explicitly during file ingestion:

- Author name
- Source URL
- Last modified timestamp
- Document file type (.pdf/.docx)

Semantic Meta

Generated dynamically by LLMs or classifiers:

- Text Sentiment score
- Automatic Topic Tagging
- Reading level difficulty
- Intended customer cohort

Ashish Jangra

[linkedin.com/in/ashish-jangra](https://www.linkedin.com/in/ashish-jangra)



CONCEPT UNLOCKED

- 📌 Prevents "lost in space" semantic search blindness
- ⚡ Single-Stage filtering evaluates index & tags simultaneously
- 🔒 Essential for high-security multi-tenant compliance
- 🧬 Combines structural metadata with AI-generated semantic tags

CONNECT WITH ME



LinkedIn



GitHub



Kaggle



Instagram

Ashish Jangra

[linkedin.com/in/ashish-jangra](https://www.linkedin.com/in/ashish-jangra)

