

TOPIC 48/60

Hybrid Search in RAG

Lexical Meets Semantic.

How merging traditional keyword search with high-dimensional vector embeddings yields unbreakable production retrieval systems.

Ashish Jangra

linkedin.com/in/ashish-jangra



THE KEYWORD TRAP

The Blindness of Keywords

Traditional keyword search (like BM25 / TF-IDF) relies on matching exact characters. While it is incredibly fast, it lacks an understanding of synonymy.

The Lexical Void

If a user queries: *"How do I cure my **migraine**?"*

A pure keyword index fails to match highly relevant documents discussing *"treatments for severe headaches"* because the literal character matches are completely absent.

Ashish Jangra

[linkedin.com/in/ashish-jangra](https://www.linkedin.com/in/ashish-jangra)



THE SEMANTIC TRAP

The Blindness of Vector Space

Conversely, dense vector embeddings are highly general conceptual maps. They can fail miserably when trying to locate specific, exact strings.

The Semantic Leak

If a developer searches for: *"Error code **CVE-2026-4011**"* or a product serial number: *"**M1-MAX-16G**"*

Vector search may group those coordinates next to generic terms like "system security flaw" or "laptop chip," failing to fetch the actual document containing the exact tag.


Ashish Jangra

[linkedin.com/in/ashish-jangra](https://www.linkedin.com/in/ashish-jangra)



THE CONVERGENCE

Parallel Pipelines



What is Hybrid Search?

Hybrid Search executes two separate queries in parallel: a **lexical (Sparse) keyword search** and a **conceptual (Dense) semantic search**, combining their output.

The Dual Retrieval Stack:

- **Sparse Match (BM25):** Fast, exact-match, rules-driven.
- **Dense Match (Embeddings):** Intent-driven, abstract similarity map.

Ashish Jangra

[linkedin.com/in/ashish-jangra](https://www.linkedin.com/in/ashish-jangra)



TAXONOMY MATRIX

Sparse vs. Dense Vectors

Dimension	Sparse Vectors	Dense Vectors
Representation	Millions of dimensions (mostly zeroes)	768 to 1536 compact float decimals
Matching Rule	Frequency & Term overlap (BM25)	Cosine Distance in coordinate space
Best For	IDs, part numbers, exact code strings	Synonyms, concepts, semantic meaning

* Note: Modern databases (e.g. Qdrant, Pinecone) can store and index BOTH sparse and dense vectors within a single collection payload.



THE SCORING NIGHTMARE

The Scoring Clash

How do we merge results from two radically different scoring structures into a single ranking list?

BM25 Scoring Range

Based on inverse document frequency.
Logarithmic curves with no absolute upper bound.

Score Range: [0 to $+\infty$]

Cosine Similarity Range

Measures vector orientation angles inside high-dimensional spatial graphs.

Score Range: [-1.0 to 1.0]

⚠ Warning: You cannot simply add a BM25 score of 14.2 to a Cosine Similarity score of 0.82. The math is apples and oranges.

Ashish Jangra

[linkedin.com/in/ashish-jangra](https://www.linkedin.com/in/ashish-jangra)



THE ALGORITHM

RRF Algorithm

Instead of relying on unstable raw scores, **Reciprocal Rank Fusion (RRF)** combines scores based solely on the **ranks** (positions) of elements.

// Reciprocal Rank Fusion Formula

$$RRF(d \in D) = \sum_{m \in M} \frac{1}{k + r_m(d)}$$

- $r_m(d)$: The position rank of document d inside retrieval list m .
- k : A smoothing constant (commonly set to 60) to prevent top ranks from completely dominating the score list.
- This algorithm scores documents consistently across disparate databases without manual normalization tuning.

Ashish Jangra

[linkedin.com/in/ashish-jangra](https://www.linkedin.com/in/ashish-jangra)



STAGE 2 VALIDATION

Two-Stage Hybrid

Stage 1: Fast Retrieval (Bi-Encoders)

Execute BM25 and vector search simultaneously. Retrieve a broad list of the top 50-100 candidates based on speed optimizations.

Stage 2: Precision Re-ranking (Cross-Encoders)

Send candidates to a slow but hyper-precise Cross-Encoder re-ranking model. It performs detailed token-to-token semantic attention to produce the final top 5 output.

Ashish Jangra


[linkedin.com/in/ashish-jangra](https://www.linkedin.com/in/ashish-jangra)



SYSTEM DESIGN

The Production Flow

How a raw user query flows through a high-performance Hybrid Search architecture:





Ashish Jangra

[linkedin.com/in/ashish-jangra](https://www.linkedin.com/in/ashish-jangra)



CONCEPT UNLOCKED



CONNECT WITH ME



Ashish Jangra

[linkedin.com/in/ashish-jangra](https://www.linkedin.com/in/ashish-jangra)

