

TOPIC 49/60

What is Reranking in RAG?

Precision meets Scale.

How a two-stage retrieval pipeline sorts vector noise to feed LLMs only the absolute gold.

Ashish Jangra

linkedin.com/in/ashish-jangra



THE PROBLEM

The Vector Search Noise

Vector search is fantastic at **Recall** (finding the top 100 documents that *might* be relevant). However, it struggles heavily with **Precision** (finding the exact 3 paragraphs that answer the specific question).

Why Vector Search Recovers Noise

When compressing millions of documents into standalone vector embeddings (e.g. 1536 float decimals), we lose structural nuances, fine syntactic rules, and delicate keyword details. The result? A query for *"Which model outperforms GPT-4 in math?"* retrieves documents that discuss *math, GPT-4, and models*, but lacks the exact mathematical comparison.

Ashish Jangra

[linkedin.com/in/ashish-jangra](https://www.linkedin.com/in/ashish-jangra)



THE ROOT CAUSE

The Encoder Clash



Bi-Encoder vs. Cross-Encoder

Embedding databases use **Bi-Encoders** because they are lightning fast, but they struggle with complex relationship comparisons. Reranking introduces a **Cross-Encoder**.

Bi-Encoders (Standard RAG)

Embeds queries and documents independently. Compares them instantly using simple cosine distance. Great for rapid filtering at scale, but misses precise attention matches.

Cross-Encoders (Reranking)

Processes the Query and Document together in a single pass. Calculates joint, word-to-word cross-attention weights. Hyper-accurate, but mathematically expensive.

Ashish Jangra

[linkedin.com/in/ashish-jangra](https://www.linkedin.com/in/ashish-jangra)



ARCHITECTURE

The 2-Stage Retrieval Stack

How elite enterprise pipelines achieve both high scale AND maximum accuracy:

Stage 1: Broad Recall (Bi-Encoder)

Search 10 million vector points in 5 milliseconds. Retrieve the top 100 generic candidates. Extremely cheap, maximum scale, average precision.

Stage 2: Micro-Precision (Cross-Encoder Reranker)

Pass ONLY those 100 candidates to a Cross-Encoder model. Reranker recalculates deep contextual scores, sorting the top 5 absolute best hits to feed the LLM.

Ashish Jangra

[linkedin.com/in/ashish-jangra](https://www.linkedin.com/in/ashish-jangra)



HOW IT WORKS

Token-to-Token Attention

Standard embeddings are static representations. A Reranker evaluates text dynamically:

```
// Cross-Encoder Input Architecture
```

```
Input = [CLS] Query [SEP] Document_Candidate [SEP]
```

This joint sequence forces every word in the Query to map attention scores directly against every word in the Document during self-attention blocks.

The model outputs a continuous probability score from [0.0 to 1.0] indicating the absolute contextual likelihood that the document answers the query.

Ashish Jangra

[linkedin.com/in/ashish-jangra](https://www.linkedin.com/in/ashish-jangra)



VISUALIZED

Sorting the Chunks

How a Reranker rearranges the original order returned by a Vector Database:

Document ID	Vector Similarity (Recall)	Reranked Score (Precision)
Doc #14	Rank 15 (Score: 0.72)	→ Rank 1 (Score: 0.96)
Doc #01	Rank 1 (Score: 0.89)	→ Rank 4 (Score: 0.42)
Doc #09	Rank 2 (Score: 0.86)	→ Rank 2 (Score: 0.81)

A document can have low physical vector similarity but extreme situational contextual relevance. Rerankers catch these hidden signals.



ECONOMICS

API & Token Optimization

Sending dozens of raw, unverified chunks directly to an LLM context window causes severe operational issues:

1. Context Bloat Costs

Every garbage chunk sent to a frontier model (like GPT-4) costs per-token fees. Compressing 100 retrieved passages down to the 3 absolute best hits reduces inference fees by up to 80%.

2. Improved Latency

Less context to read means faster Time-to-First-Token (TTFT) during streaming.

3. Eliminating "Lost in the Middle"

Fewer documents in the prompt reduces the risk of the LLM hallucinating or overlooking the exact fact it needs to formulate the response.

Ashish Jangra

[linkedin.com/in/ashish-jangra](https://www.linkedin.com/in/ashish-jangra)



THE LANDSCAPE

The Reranking Stack

Proprietary APIs (Cohere Rerank)

Simple, cloud-hosted endpoints. Excellent multi-lingual support, zero training required, plug-and-play optimization.

Open-Source Giants (BGE / Jina Rerank)

Run locally on your own GPUs. Highly customizable, supports domain-specific training (SFT), and fully private.

Late Interaction (ColBERT)

A middle ground that scales cross-attention speed by indexing separate token-level vectors, bypassing full cross-encoder cost thresholds.

Ashish Jangra


[linkedin.com/in/ashish-jangra](https://www.linkedin.com/in/ashish-jangra)



PRODUCTION FLOW

The Production Map

How data flows from query to response inside an optimized, enterprise-grade RAG pipeline:



Ashish Jangra

[linkedin.com/in/ashish-jangra](https://www.linkedin.com/in/ashish-jangra)



Reranking in RAG

CONCEPT UNLOCKED



Resolves the precision limitations of vector search



Cross-Encoders evaluate full context with mutual attention



Minimizes LLM prompt size and cuts VRAM/API token cost



Crucial Stage-2 mechanism in modern enterprise search

CONNECT WITH ME



LinkedIn



GitHub



Kaggle



Instagram

Ashish Jangra

linkedin.com/in/ashish-jangra

